

Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition.

Every Grader function has to return True.

Importing packages

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model

Creating custom dataset
```

```
In [2]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                           n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_classification.html) for more details

In [3]: X.shape, y.shape

Out[3]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
In [4]: #please don't change random state
# you need not standardize the data as it is already standardized
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)

In [5]: X_train.shape, y_train.shape, X_test.shape, y_test.shape

Out[5]: ((37500, 15), (37500,)), ((12500, 15), (12500,))
```

SGD classifier

```
In [6]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html)
```

```
Out[6]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                    random_state=15, verbose=2)
```

```
In [7]: clf.fit(X=X_train, y=y_train) # fitting our model

--> Epoch 1
Norm: 0.77, MWZ: 15, Bias: -0.316053, T: 37500, Avg. loss: 0.455552
Total training time: 0.02 seconds.
--> Epoch 2
Norm: 0.91, MWZ: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394688
Total training time: 0.03 seconds.
--> Epoch 3
Norm: 0.98, MWZ: 15, Bias: -0.500082, T: 112500, Avg. loss: 0.385711
Total training time: 0.04 seconds.
--> Epoch 4
Norm: 1.02, MWZ: 15, Bias: -0.650292, T: 150000, Avg. loss: 0.382083
Total training time: 0.06 seconds.
--> Epoch 5
Norm: 1.04, MWZ: 15, Bias: -0.719520, T: 187500, Avg. loss: 0.380486
Total training time: 0.07 seconds.
--> Epoch 6
Norm: 1.05, MWZ: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.08 seconds.
--> Epoch 7
Norm: 1.06, MWZ: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.09 seconds.
--> Epoch 8
Norm: 1.06, MWZ: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.10 seconds.
--> Epoch 9
Norm: 1.07, MWZ: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.11 seconds.
--> Epoch 10
Norm: 1.08, MWZ: 15, Bias: -0.853130, T: 375000, Avg. loss: 0.378630
Total training time: 0.12 seconds.
Convergence after 10 epochs took 0.12 seconds
```

```
Out[7]: SGDClassifier(eta0=0.0001, learning_rate='constant', loss='log',
                    random_state=15, verbose=2)
```

```
In [8]: clf.coef_, clf.coef_.shape, clf.intercept_
#clf.coef_ will return the weights
#clf.coef_.shape will return the shape of weights
#clf.intercept_ will return the intercept term
```

```
Out[8]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
  0.58016579, -0.45242483, -0.09408013,  0.2092732 ,  0.18684126,
  0.18705191,  0.00421515, -0.0796637 ,  0.33852802,  0.02266721]]),
(1, 15),
array([-0.8531383]))
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

- We will be giving you some functions, please write code in that functions only.
- After every function, we will be giving you expected output, please make sure that you get that output.

- Initialize the weight_vector and intercept term to zeros (Write your code in `def initialize_weights()`)
- Create a loss function (Write your code in `def logloss()`)
$$\text{logloss} = -1 * \frac{1}{n} \sum_{for each Y_i, Y_{pred}} (Y_i \log 10(Y_{pred}) + (1 - Y_i) \log 10(1 - Y_{pred}))$$
- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector (write your code in `def gradient_dw()`)
$$dw^{(i)} = x_n(y_n - \sigma((w^{(i)})^T x_n + b')) - \frac{1}{n} w^{(i)}$$
 - Calculate the gradient of the intercept (write your code in `def gradient_db()`) check this
$$db^{(i)} = y_n - \sigma((w^{(i)})^T x_n + b'))$$
 - Update weights and intercept (check the equation number 32 in the above mentioned pdf):
$$w^{(i+1)} \leftarrow w^{(i)} + \alpha(dw^{(i)})$$

$$b^{(i+1)} \leftarrow b^{(i)} + \alpha(db^{(i)})$$
 - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
 - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
 - append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)

Initialize weights

```
In [9]: def initialize_weights(row_vector):
''' In this function, we will initialize our weights and bias'''
#initialize weights as 1d array consisting of all zeros similar to the dimensions of row_vector
#you use zeros_like function to initialize zero, check this link https://docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
#initialize bias to zero
w=np.zeros_like(row_vector)
b=0
return w,b
```

```
In [10]: row_vector=X_train[10]
w,b = initialize_weights(row_vector)
print('w =',w)
print('b =',str(b))

w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0

Grader function - 1
```

```
In [11]: dim=X_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
return True
grader_weights(w,b)
```

```
Out[11]: True

Compute sigmoid
sigmoid(z) = 1/(1 + exp(-z))
```

```
In [12]: import math as m
import numpy as np
def sigmoid(z):
''' In this function, we will return sigmoid of z'''
# compute sigmoid(z) and return
return (1/(1+m.exp(-z)))

Grader function - 2
```

```
In [13]: def grader_sigmoid(z):
val=sigmoid(z)
assert(val==0.8807970779778823)
return True
grader_sigmoid(2)
```

```
Out[13]: True

Compute loss
logloss = -1 * 1/n \sum_{for each Y_i, Y_{pred}} (Y_i \log 10(Y_{pred}) + (1 - Y_i) \log 10(1 - Y_{pred}))
```

```
In [14]: def logloss(y_true,y_pred):
# you have been given two arrays y_true and y_pred and you have to calculate the logloss
#while dealing with numpy arrays you can use vectorized operations for quicker calculations as compared to using loops
#https://www.pythonlikeyoumeanit.com/Module3_IntroducingNumpy/VectorizedOperations.html
#https://www.geeksforgeeks.org/vectorized-operations-in-numpy/
#write your code here
losses=0
for i in range(len(y_pred)):-
    losses=(y_true[i]*np.log10(y_pred[i]) + ((1-y_true[i])*np.log10(1-y_pred[i]))) +losses
return ((-1*losses)/len(y_pred))

Grader function - 3
```

```
In [15]: #round off the value to 8 values
def grader_logloss(true,pred):
loss=logloss(true,pred)
assert(np.round(loss,8)==0.076449)
return True
true=np.array([1,1,0,1,0])
pred=np.array([0.9,0.8,0.1,0.8,0.2])
grader_logloss(true,pred)
```

```
Out[15]: True

Compute gradient w.r to w
dw^{(i)} = x_n(y_n - \sigma((w^{(i)})^T x_n + b')) - \frac{1}{N} w^{(i)}
```

```
In [16]: def multiplication_fun(X,Y):
sum=0
for i in range(len(X)):
sum=sum+X[i]*Y[i]
return sum1
```

```
In [17]: #make sure that the sigmoid function returns a scalar value, you can use dot function operation
def gradient_dw(x,y,w,b,alpha,N):
''' In this function, we will compute the gradient w.r to w'''
dw=x*(y-sigmoid(multiplication_fun(w,x)+b)-(alpha/N)*w)
return dw

Grader function - 4
```

```
In [18]: def grader_dw(x,y,w,b,alpha,N):
grad_dw=gradient_dw(x,y,w,b,alpha,N)
assert(np.round(np.sum(grad_dw),5)==4.75684)
return True
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.07945546, -1.14783286,
-2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
 3.67152472,  0.01451875,  2.01662888,  0.07373904, -5.54586092])
grad_y=0
grad_w=np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
-0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
-0.05708087, -0.064768 ,  0.18012332, -0.16880843, -0.27079877])
grad_b=0.5
alpha=0.0001
N=len(X_train)
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

```
Out[18]: True

Compute gradient w.r to b
db^{(i)} = y_n - \sigma((w^{(i)})^T x_n + b)
```

```
In [19]: #sb should be a scalar value
def gradient_db(x,y,w,b):
''' In this function, we will compute the gradient w.r to b'''
db=y-sigmoid(multiplication_fun(w,x)+b)
return db

Grader function - 5
```

```
In [20]: def grader_db(x,y,w,b):
grad_db=gradient_db(x,y,w,b)
assert(np.round(grad_db,4)==-0.3714)
return True
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.07945546, -1.14783286,
-2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
 3.67152472,  0.01451875,  2.01662888,  0.07373904, -5.54586092])
grad_y=0.5
grad_b=0.1
grad_w=np.array([ 0.03364887,  0.03612727,  0.02786927,  0.08547455, -0.12870234,
-0.02555288,  0.11858013,  0.13305576,  0.07310204,  0.15149245,
-0.05708087, -0.064768 ,  0.18012332, -0.16880843, -0.27079877])
alpha=0.0001
N=len(X_train)
grader_db(grad_x,grad_y,grad_w,grad_b,)
```

```
Out[20]: True

# prediction function used to compute predicted_y given the dataset X
def pred(w,b,X):
N = len(X)
predict = []
for i in range(N):
z=np.dot(w,X[i])+b
predict.append(sigmoid(z))
return np.array(predict)
```

Implementing logistic regression

```
In [22]: import math
def train(X_train,y_train,X_test,y_test,epochs,alpha,eta,N):

# Initialize the weights
#write your code to perform SGD
# iterating over each epochs
#for each instance of the table
w,b=initialize_weights(X_train[0]) #weight initialization
train_loss_list,test_loss_list=[],[] # creating the list which store the train & test score
for i in range(0,epochs): #for each epoch
for i in range(0,len(X_train)): #for each data point
w1=gradient_dw(X_train[i],y_train[i],w,b,alpha,N)
b1=gradient_db(X_train[i],y_train[i],w,b)
w=w+(eta*w1) #updating the weight
b=b+(eta*b1) #updating the intercept
y_pred_train=pred(w,b,X_train)
training_loss=logloss(y_train,y_pred_train)
train_loss_list.append(training_loss)
y_pred_test=pred(w,b,X_test)
testing_loss=logloss(y_test,y_pred_test)
test_loss_list.append(testing_loss)
return w,b,train_loss_list,test_loss_list
```

```
In [23]: alpha=0.0001
eta0=0.0001
N=len(X_train)
epochs=20
w,b,train_loss,test_loss=train(X_train,y_train,X_test,y_test,epochs,alpha,eta0,N)
```

```
In [24]: import numpy as np
```

```
In [25]: #print thr value of weights w and bias b
print(w)
print(b)

[-4.20394719e-01  1.02911530e-01 -1.48319219e-01  3.38095823e-01
-2.20731204e-01  5.69659877e-01 -4.45186059e-01 -9.00099047e-02
2.21508210e-01  1.73585090e-01  1.98838397e-01 -4.13161582e-04
-8.1124978e-02  3.38070557e-01  2.29309039e-02
-0.8897519366656894]
```

```
In [26]: # these are the results we got after we implemented sgd and found the optimal weights and intercept
w-clf.coef_, b-clf.intercept_

array([[ -0.0660278 ,  0.00743588,  0.00027114, -0.00334825, -0.0125445 ,
 0.00723077,  0.00407822,  0.01232502, -0.00725325,
 0.00148649, -0.00463232, -0.0015213 ,  0.00064254,  0.00628959]],
array([-0.03661364]))
```

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in order of 10^-2

Grader function - 6

```
In [27]: #this grader function should return True
#the difference between custom weights and clf.coef_ should be less than or equal to 0.05
def difference_check_grader(w,b,coef_,intercept):
val_array=np.abs(np.array(w-coef_))
assert(np.all(val_array<=0.05))
print('The custom weights are correct')
return True
difference_check_grader(w,b,clf.coef_,clf.intercept_)
```

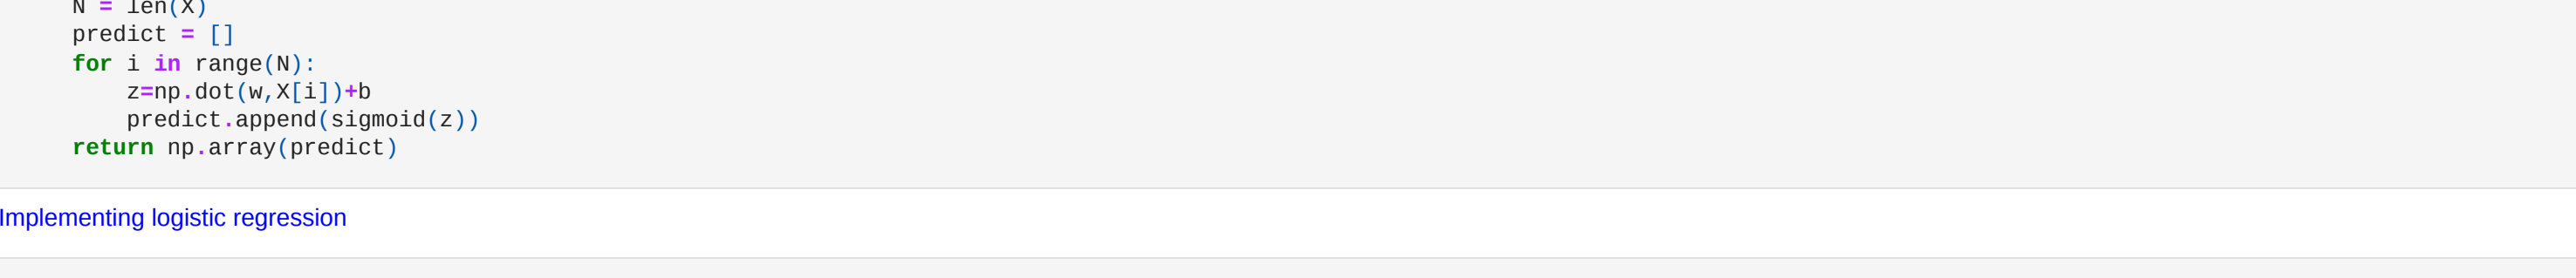
The custom weights are correct

```
Out[27]: True

Plot your train and test loss vs epochs
plot epoch number on X-axis and loss on Y-axis and make sure that the curve is converging
```

```
In [28]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
In [29]: plt.plot(train_loss, 'g', label='Training_Loss')
plt.plot(test_loss, 'r', label='Testing_Loss')
plt.title('loss changes on increasing epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [ ]:
```

```
In [30]: print(len(train_loss),len(test_loss))

20 20
```